

Visual Computing Fundamentals

Image Processing

Assignment 2

Adrian Opheim

November 10, 2017

Task 1: Theory Question

- a) The convolution theorem can be seen in Equation 1, where \mathcal{F} is the Fourier transform, $*$ is the convolution operator and \cdot is pointwise multiplication.

i) *What does the convolution theorem entail?*

The convolution theorem entails that convolution is commutative, because the same result would be obtained if the order of f and g were reversed. The convolution theorem also states that convolution in the frequency domain is analogous to multiplication in the spatial domain. The two are being related by the forward and inverse Fourier transforms, respectively.

$$\mathcal{F}\{f * g\} = \mathcal{F}\{f\} \cdot \mathcal{F}\{g\} \quad (1)$$

ii) *What are the main steps an implementation of frequency filtering requires?*

The basic filtering equation has the form:

$$g(x, y) = \mathcal{F}^{-1} [H(u, v) \cdot F(u, v)] \quad (2)$$

where \mathcal{F}^{-1} is the Inverse Discrete Fourier Transform (IDFT), $F(u, v)$ is the Discrete Fourier Transform (DFT) of the input image $f(x, y)$. $H(u, v)$ is a filter function and $g(x, y)$ is the filtered output image. A step-by-step procedure of filtering in the frequency domain is given here: ¹

- (i) Given an input image of size $M \times N$, obtain the padding parameters P and Q from $P \geq 2M - 1$ and $Q \geq 2N - 1$. Typically, we select $P = 2M$ and $Q = 2N$.
- (ii) Form a padded image, $f_p(x, y)$ of size $P \times Q$ by appending the necessary number of zeros to $f(x, y)$.
- (iii) Multiply $f_p(x, y)$ by $(-1)^{x+y}$ to center its transform.
- (iv) Compute the DFT, $F(u, v)$, of the image from step 3.
- (v) Generate a real, symmetric filter function, $H(u, v)$ of size $P \times Q$ with center at coordinates $(P/2, Q/2)$. Form the product $G(u, v) = H(u, v)F(u, v)$ using array multiplication; that is, $G(i, k) = H(i, k)F(i, k)$
- (vi) Obtain the processed image:

$$g_p(x, y) = \{\text{real} [\mathcal{F}^{-1} [G(u, v)]]\} (-1)^{x+y} \quad (3)$$

¹Taken from page 285 of *Digital Image processing*

where the real part is selected in order to ignore parasitic complex components resulting from computational inaccuracies, and the subscript p indicates that we are dealing with padded arrays.

- (vii) Obtain the final processed result $g(x, y)$ by extracting the $M \times N$ region from the top, left quadrant of $g_p(x, y)$.
- b) Convolution kernels are typically understood in terms of the frequency domain. In your own words, explain what high- and low-pass filters are.

A low-pass filter is a filter that discards high frequencies obtained from the DFT, while passing low frequencies. A low-pass filter blurs an image.

A high-pass filter has the opposite property: It discards low frequencies and passes high frequencies. A high-pass filter would therefore enhance sharp detail, but cause a reduction in contrast in the image.

- c) The Fourier spectrum $|\mathcal{F}\{f\}|$ of three commonly used convolution kernels can be seen in Figure 1. For each kernel, write down what kind of kernel it is. Explain your reasoning.

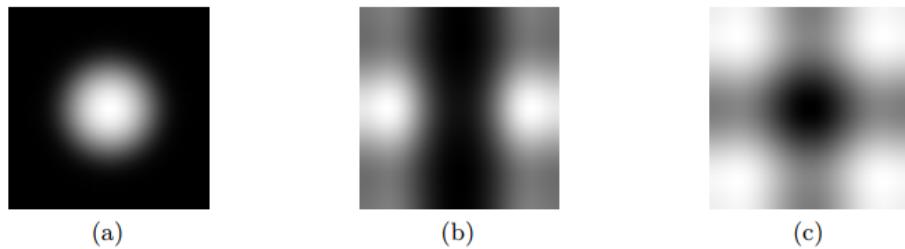


Figure 1: Three different convolution kernels

Kernel (a) is a low-pass filter, probably a Gaussian Lowpass-filter, as we have no reoccurring white circles.

I would assume (b) and (c) are both high-pass filters, because they are both rejecting low frequencies in the middle, and passes high frequencies further away from the origo.

- d) Answer the following:
 - i) *Explain why padding is important when we want to do frequency filtering.*
If the images are not padded, we could expect a wraparound error. Meaning that the filtering will have no effect, as the filter will meet light areas in both the image and in its right neighbour.
 - ii) *In the context of frequency filtering, we commonly pad images by appending zeros to the rows and columns as seen in Figure 2. Let's assume we instead pad the images by prepending zeros to the rows and columns. This can be seen in Figure 2 (b). Let the number of zeros we pad by remain the same. Does this change the result of frequency filtering? Explain your reasoning.*
This will not change the result of frequency filtering, as you can avoid the wraparound error by appending or prepending zeros to the image. The number of zeros should satisfy $P \geq A + B - 1$, where P is the total length of functions f and h , composed of A and B samples. If each function has 400 points, so the minimum value that can be used is $P = 799$, which implies that we could append or prepend 399 zeros to each function.

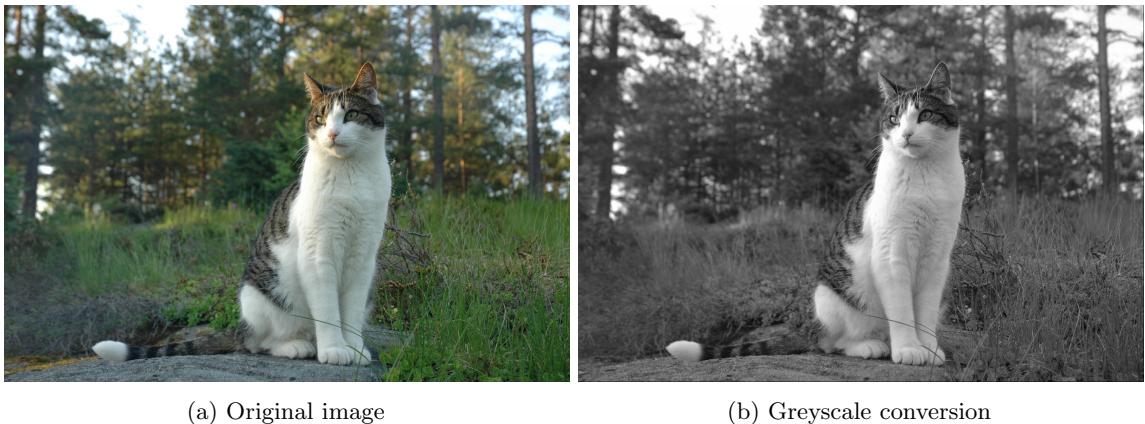


Figure 2: (a) shows appending zeros, (b) shows prepending zeros.

Task 2: Frequency Domain Filtering

- a) *Using an already existing FFT implementation, implement a function that takes a greyscale image and an arbitrary linear convolution kernel and performs frequency filtering. Apply the function you just made on an appropriate greyscale image with both a high- and a low-pass convolution kernel. It's up to you to choose which convolution kernels to use.*

The chosen image and the following greyscale conversion, using the luminance-preserving method used in Assignment 1, is shown in Figure 3



(a) Original image

(b) Greyscale conversion

Figure 3: Original image and corresponding greyscale conversion. The image is of size (997×1500) pixels

Step (i) Obtaining padding parameters

As the grayscale image is of dimensions $(M \times N) = (997 \times 1500)$, the padding parameters are set to $P = 2M = 1994$ and $Q = 2N = 3000$.

Step (ii) Padding the image

Forming a padded image of size $P \times Q$, by appending zeros. The corresponding image can be seen in Figure 4



Figure 4: The padded image of size 1994×3000

Step (iii) Centering the image's transform

After centering the image $f_p(x, y)$ with $(-1)^{x+y}$, we get the image shown in Figure 5

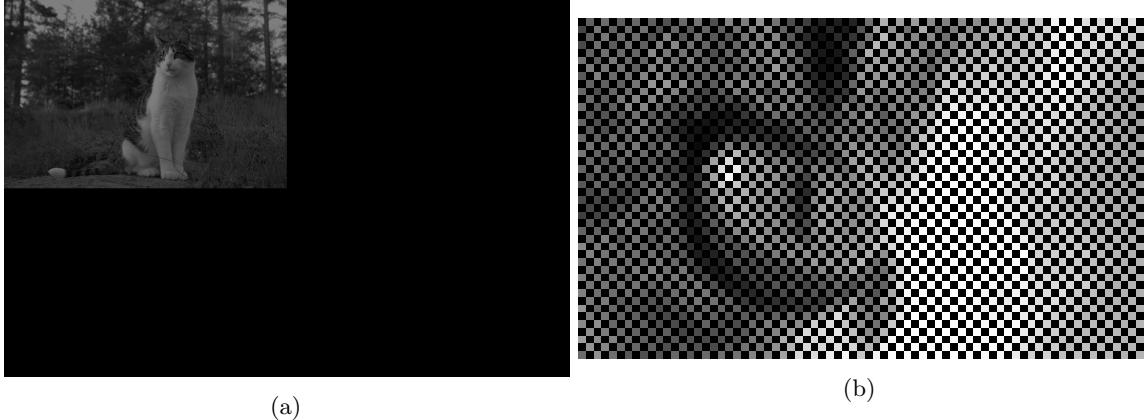


Figure 5: (a) shows the complete, centered image. (b) shows an even further zoom, making it possible to see individual pixels that are changed to alternating black, and original grayscale value. The black pixels are pixels set to a value of zero (black), for visualisation purposes, as the pixel's actual value is the negative of the original pixel value.

Step (iv) Computing the Fourier transform The Fourier transform of the image can be seen in Figure 6

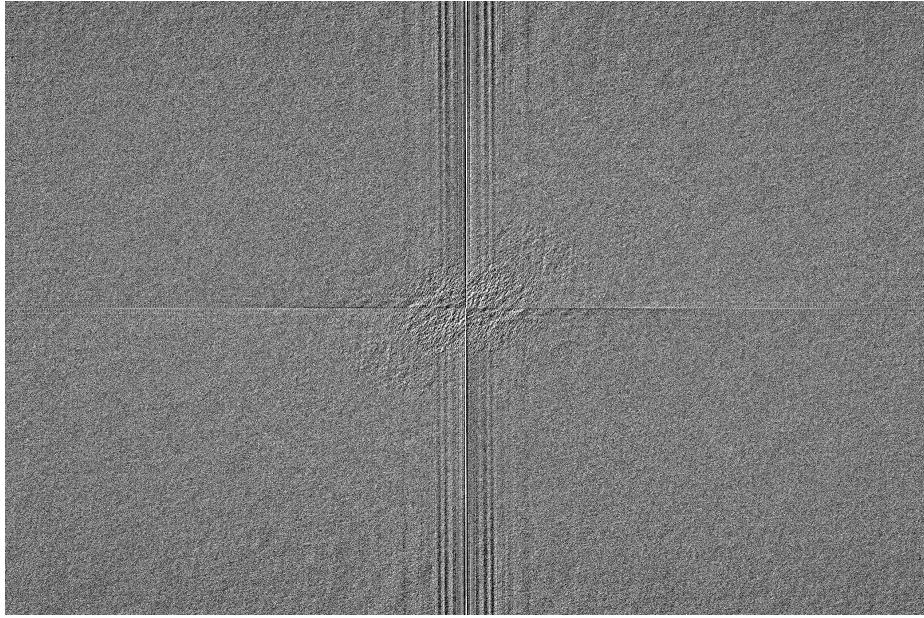


Figure 6: The Fourier transform of the padded image.

Step (v) Generating a filter

A Gaussian lowpass filter of size $(P \times Q)$ is generated. The parameters used are $D_0 = 100$. The generated lowpass filter can be seen in Figure 7 (a), and the product $G(u, v) = H(u, v)F(u, v)$ can be seen in Figure 7 (b).

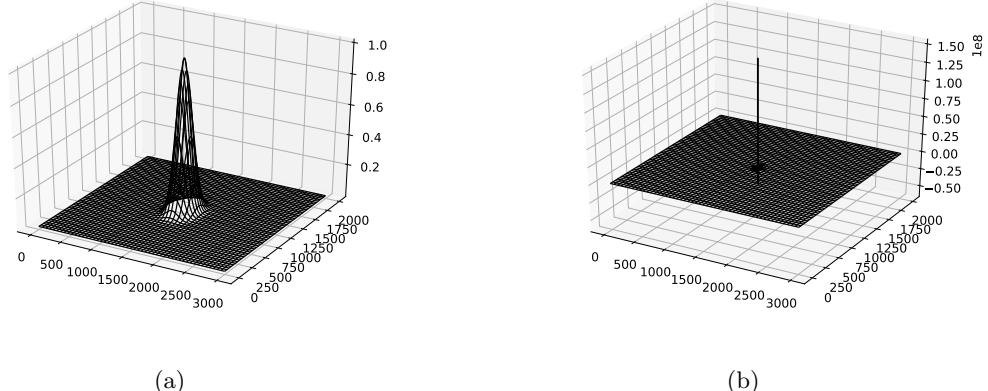


Figure 7: (a) shows the generated Gauss lowpass kernel. (b) shows the product $G(u, v) = H(u, v)F(u, v)$.

Step (vi) Obtaining the processed image

The processed image can be seen in Figure 8



Figure 8: The inverse Fourier transform of the padded image.

Step (vii) Obtaining the image of size $(M \times N)$

The filtered image can be seen in Figure 9



Figure 9: (a) shows the processed image after low-pass filtering, while (b) shows the original grayscale image.

Applying a high-pass filter

A Butterworth high-pass filter is also implemented. The high-pass filter is useful for sharpening images. The generated high-pass Butterworth filter is shown in Figure 12

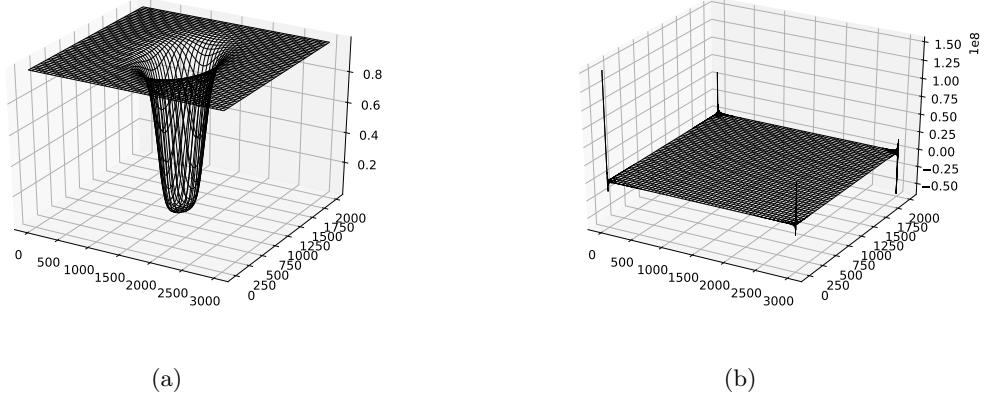


Figure 10: (a) shows the generated Butterworth filter with $D_0 = 300, n = 3$ (b) shows the spectrum of the Fourier transformed image, $F(x, y)$. Note that this spectrum is not centered, as I did not get the high-pass filter to work with the centered image.

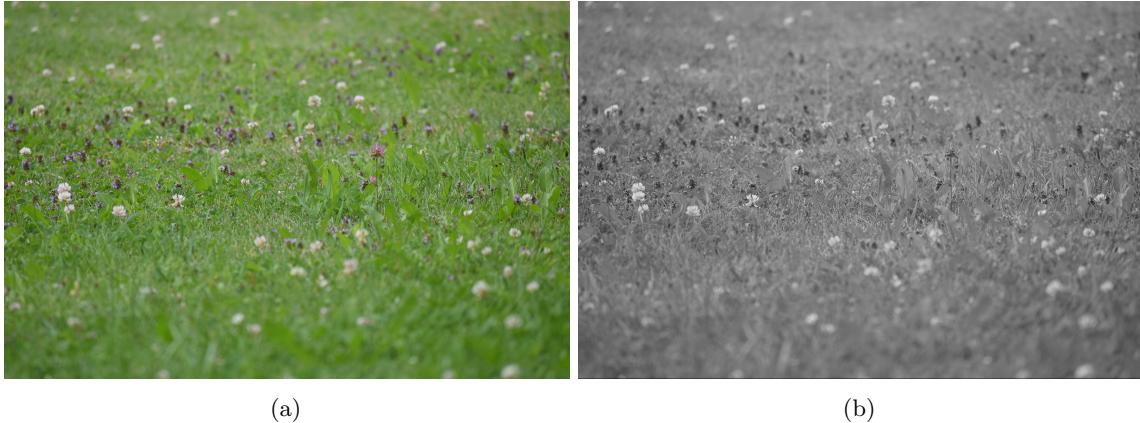
Figure 11 shows the processed image after the high-pass filtering, compared to the original grayscale image.



Figure 11: (a) shows the processed image with a high-pass Butterworth filter, while (b) shows the original grayscale image.

Task 3: Unsharp Masking

- Implement a function that performs unsharp masking on a greyscale image using frequency filtering. Instead of implementing Equation 3 directly you must compute a single kernel, which when used with frequency filtering will yield the desired sharpened image. We recommend that you use Gaussian smoothing when creating the unsharp kernel.*



(a)

(b)

Figure 12: (a) shows the original color image used for unsharp masking. (b) shows the corresponding grayscale image, generated using the luminance-preserving method.

I did not completely understand what was being meant by "computing a single kernel". But I have implemented the procedure in the book, going as follows:

$$g(x, y) = f(x, y) + k \cdot g_{\text{mask}}(x, y)$$

where

$$g_{\text{mask}} = f(x, y) - f_{LP}(x, y)$$

and

$$f_{LP}(x, y) = \mathcal{F}^{-1} [H_{LP}(u, v)F(u, v)]$$

... and this was how far I got before the deadline, unfortunately. Very much time was spent fiddling with NumPy and with exporting figures. I will try to use my time smart for the next assignment ;)

For future reference, my current code is added below. It is very much unpolished.

```

1 import numpy as np      #n-D arrays
2 import matplotlib.pyplot as plt          # For plotting
3 from mpl_toolkits.mplot3d import axes3d # For 3D plotting
4 from PIL import Image
5 import math
6
7
8 def img_save(img, filename):
9     print("Saving ...")
10    img.save(filename, "PNG")
11
12 def img_show(img):
13     plt.figure()
14     plt.imshow(img, cmap='gray')      # have to use the gray color map.
15     plt.axis('off')
16     plt.show()
17
18
19 def conv2Grey_lum(img):
20     img_arr = np.array(img) # convert from image object to array
21     r = img_arr[:, :, 0] # equivalent to img[:, :, 0]

```

```

22     g = img_arr[:, :, 1]
23     b = img_arr[:, :, 2]
24
25     grey = np.zeros((len(r), len(r[0])), np.uint8)
26     print("Converting to grey...")
27     for row in range(0, len(r)-1):
28         for col in range(0, len(r[0])-1):
29             grey[row][col] = int(0.2126*r[row][col] + 0.7152*g[row][col]
30                         + 0.0722*b[row][col])
31
32     img = Image.fromarray(grey) # converting back to image object
33     return img
34
35
36 def freq_filter(img):
37
38     def pad_image(img_arr):
39         print("Padding image...")
40         img_padded = np.zeros((P, Q), dtype = np.uint8)
41         for i in range(0, M):
42             for j in range(0, N):
43                 img_padded[i][j] = img_arr[i][j]
44         return img_padded
45
46     def plot_wireframe(kernel, filename):
47
48         u = np.zeros((P, Q), dtype=np.int)
49         v = np.zeros((P, Q), dtype=np.int)
50         j_counter = 0
51         for i in range(0, P):
52             for j in range(0, Q):
53                 u[i] = np.arange(Q)
54                 v[i].fill(j_counter)
55                 j_counter += 1
56
57         #print("u: ", u, "\nv: ", v, "\nkernel: ", kernel)
58         #print("u: ", u.shape, "v: ", v.shape, "kernel: ", kernel.shape)
59         fig = plt.figure()
60         ax = fig.add_subplot(111, projection='3d')
61         lines = ax.plot_wireframe(u, v, kernel)
62         plt.setp(lines, color='black', linewidth=0.5)
63         plt.show()
64         fig.savefig(filename, format='pdf') #saving as pdf
65
66
67 def D(u, v): # Function for calculating the distance between a point
68     # (u, v) in the frequency domain and the center of the frequency
69     # rectangle
70     term1 = (u - P/2)*(u - P/2)
71     term2 = (v - Q/2)*(v - Q/2)
72     #print(math.sqrt(term1 + term2))
73     ans = math.sqrt(term1 + term2)

```

```

72     # Not possible to divide by zero:
73     if (ans < 10e-7):
74         return 10e-7
75     else:
76         return ans
77
78 def lowpass_gauss(D0):
79     H = np.zeros((P, Q), dtype=np.float) # H is the kernel —> H(u,
80                 v)
81     print("Creating Gauss kernel...")
82     for u in range(0, P):
83         for v in range(0, Q):
84             teller = -(D(u, v) * D(u, v))
85             nevner = 2 * (D0 * D0)
86             eksp = teller / nevner
87             H[u][v] = math.exp(eksp)
88
89     #print(H.shape)
90     #plot_wireframe(H)
91     #print("H: ", H)
92     #img = Image.fromarray(np.real(H))
93     #img = img.convert('RGB')
94     #img.save(img, "./img-proc/lowpass-gauss.png")
95     return H
96
97 def highpass_butterworth(D0, n):
98     H = np.zeros((P, Q), dtype=np.float) # H is the kernel —> H(u,
99                 v)
100    print("Creating high-pass Butterworth kernel...")
101    for u in range(0, P):
102        for v in range(0, Q):
103            teller = 1
104            nevner = D0 / (D(u, v))
105            nevner = math.pow(nevner, 2*n)
106            nevner = 1 + nevner
107            H[u][v] = teller / nevner
108
109    plot_wireframe(H, './plots/highpass_butterworth.pdf')
110    return H
111
112
113
114
115 def center_transform(img):
116     print("Centering the image...")
117     #print("Image before centering: ", img)
118     centered = np.zeros((P, Q), dtype=np.float) # Have to allow
119                 negative values
120     for x in range(0, P):
121         for y in range(0, Q):
122             eksp = x + y

```

```

122     centered[x][y] = img[x][y] * math.pow(-1, eksp)
123 #print("Image after centering: ", centered)
124     return centered
125
126 def inv_fourier(G):
127     print("Inverse Fourier transform...")
128     g_p = np.fft.ifft2(G) # performing inverse Fourier transform...
129     g_p = np.real(g_p) # extracting the real part...
130     g_p = center_transform(g_p) # Centering
131     g_p = g_p.astype(np.uint8) # Converting back to uint8 type
132     return g_p
133
134 def plot_fourier(F):
135     print(F)
136     # Extracting the real part of the transform:
137     F = np.fft.fftshift(F)
138     F_log = np.log2(1 + np.abs(F)**2)
139
140     print(F_log)
141     print("max: ", np.max(F_log))
142
143     img = Image.fromarray(F_log)
144     img = img.convert('RGB')
145     img.show(img)
146
147
148
149
150
151     img_arr = np.array(img)
152     img_arr = img_arr[:, :, 0] # Removing color bands that are still
153     there.
154
155     # Getting image parameters: Original image of size (M, N)
156     M = len(img_arr)
157     N = len(img_arr[0])
158     # Size of the filter kernel:
159     P = 2*M
160     Q = 2*N
161     print("Dimensions of the image: \nM: ", M, "N: ", N, "P: ", P, "Q: "
162           , Q)
163
164     # Padding the image:
165     img_padded = pad_image(img_arr)
166
167     # Centering its transform:
168     img_centered = center_transform(img_padded)
169
170
171     # Computing the Fourier transform:
172     print("Computing Fourier transform...")
173     img_fourier = np.fft.fft2(img_padded)

```

```

173 plot_fourier(img_fourier)
174
175
176 # Creating a kernel:
177 kernel = lowpass_gauss(1e2)
178 #kernel = highpass_butterworth(3e2, 3)
179
180 # Multiplying element-wise kernel with Fourier transformed image:
181 print("Multiplying kernel with Fourier image...")
182 img_G = np.multiply(kernel, img_fourier)
183 plot_fourier(img_G)
184
185
186 # Obtaining the finished processed image:
187 img_proc = inv_fourier(img_G)
188
189
190 # Clipping out the original (M, N) sized picture:
191 img_proc = img_proc[0:M, 0:N]
192 img = Image.fromarray(img_proc)
193 img = img.convert('RGB')
194 img.show(img)
195
196 # Subtracting the low-pass image from the original grayscale:
197 g_mask = np.subtract(img_arr, img_proc)
198
199
200 # Multiplying the mask with a k value
201 g_mask = np.multiply(1, g_mask)
202
203 g = np.add(img_arr, g_mask)
204 img = Image.fromarray(g)
205 img.save(img, "./img_proc_2/img_g_k1.png")
206
207
208 return img
209
210
211 def main():
212     ,
213
214     # Creating the greyscale image:
215     img = Image.open("./images/DSC_0784.JPG")
216     img_grey = conv2Grey_lum(img)
217     img_grey.save(img_grey, "./img_proc_2/img_grey.png")
218
219     ,
220
221     img = Image.open("./img_proc/img_sushi_grey_mini.png")
222     #img = Image.open("./img_proc_2/img_grey_mini.png")
223     #img = Image.open("./images/noise-a.tiff")
224
225
226     ## Filtering in the frequency domain: ##
227     img_freq_filter = freq_filter(img)

```

```
226  
227  
228  
229  
230  
231 if __name__ == "__main__":  
232     main()
```